KENT STATE UNIVERSITY

ADVANCE MACHINE LEARNING
FINAL REPORT

DEEP LEARNING FOR BRAIN TO TEXT DECODING

INSTRUCTOR
CHAOJIANG (CJ) WU

SUBMITTED BY:
VIDUSHI VILLYAL

DECEMBER 4, 2025

1. INTRODUCTION

The topic of brain-to-text decoding has evolved from a theoretical concept to a crucial focus within contemporary artificial intelligence. Fundamentally, this technology tries to translate the electrical and magnetic activity of the brain into comprehensible writing. These brain impulses encapsulate an individual's thoughts, mental speech, planned expressions, or motor orders connected to speaking. It is feasible to deduce someone's verbal goals even when they are unable to express them physically by using deep learning algorithms to analyse these signals.

The societal implications of this innovation are substantial. Millions of people worldwide suffer from diseases like ALS, spinal cord injuries, brainstem strokes, and severe paralysis that impair their speech. Conventional assistive communication devices, such as switch-based systems or eye-tracking keyboards, frequently produce slow interaction speeds that don't match the natural rhythm of human thought. In contrast, deep learning-driven neural decoding presents a remarkable opportunity to restore conversational capabilities, autonomy, and emotional expression for affected individuals

Furthermore, brain-to-text decoding represents a new stage in understanding language processing in the brain. Training neural networks on brain signals provides researchers with valuable insights into the neural frameworks governing speech production, semantic understanding, and mental imagery. Because of this twofold benefit-medical empowerment combined with scientific investigation-brain-to-text decoding is currently one of the most fascinating fields in cutting-edge machine learning research.

2. STATE-OF-THE-ART DEEP LEARNING MODELS FOR BRAIN-TO-TEXT DECODING

Because brain signals are noisy and dynamic, translating language from neural activity is a complex task. By directly learning from raw data and recognising patterns that match the neural mechanisms underlying speech intention, deep learning models are especially skilled at handling this complexity. Over the last decade, various deep learning architectures-CNNs, RNNs, Transformers, and generative models-have uniquely contributed to advancing brain decoding methodologies.

2.1 CNNS, OR CONVOLUTIONAL NEURAL NETWORKS

CNNs frequently serve as foundational components in brain decoding pipelines due to their proficiency in recognising spatial patterns. Whether derived from ECoG grids or EEG electrode arrays, brain signals showcase significant spatial relationships where adjacent electrodes capture correlated neural activities.

Contributions of CNNs:

- They traverse across channels to identify speech-related patterns such as specific frequency bursts associated with vowel or consonant imagination.
- They pinpoint cortical regions engaged during tongue, lip, and jaw movements.
- They condense raw neural information into relevant feature maps interpretable by subsequent models.

Significance of CNNs:

Although brain signals fluctuate rapidly, CNNs can consistently identify recurring neural signatures linked to phonetic units and articulatory muscle representations.

2.2 GRUS, LSTMS, AND RECURRENT NEURAL NETWORKS (RNNS)

Since language develops over time, the brain activates corresponding neural patterns that reflect phonetic transitions and cognitive planning when a person imagines or intends to speak.

- RNNs are important because they maintain memory over long neural sequences.
- Long-range dependencies that are crucial for deciphering multi-word sentences are efficiently captured by LSTMs.

- GRUs facilitate quicker training suitable for real-time applications.

They are particularly good at decoding the following:
Neural indicators indicate phoneme transitions, imagined sequences, and temporal rhythms of speech.

Successful continuous sentence decoders were first developed using RNN-based systems.

2.3 TRANSFORMERS AND ATTENTION MECHANISMS

Transformers exemplify the current pinnacle in brain decoding technology.

Reasons for Transformer superiority over prior models:

- They analyse entire sequences concurrently rather than sequentially.
- Attention layers emphasise crucial temporal points within neural data.
- Their scalability accommodates extensive datasets efficiently.
- They manage lengthy sentences without losing contextual integrity.

Transformers have enabled:

- Semantic-level interpretation (decoding meaning beyond mere words).
- Rapid communication rates (Stanford achieved 62 words per minute).
- Non-invasive methods utilising MEG through attention-based modelling techniques.

Transformers currently represent the most robust architecture within research focused on brain-to-text translation.

2.4 ENCODER-DECODER ARCHITECTURES

These systems integrate multiple deep learning frameworks into cohesive workflows:
Encoder: Extracts significant features from neural input
Decoder: Generates coherent words or phrases based on those features

This framework parallels machine translation (e.g., translating French into English), but here it translates from neural inputs into text outputs.
Applications include:

- Real-time generation of sentences
- Decoding imagined speech
- Interpreting motor intentions related to articulation

2.5 GENERATIVE MODELS (GANS, VAES, DIFFUSION MODELS)

Generative models represent recent advancements within this domain by:

- Synthesising artificial data sets from limited samples
- Reconstructing spoken audio based on identified features
- Amending gaps found in partial signal recordings

While still experimental in nature, these models possess significant potential moving forward.

3. LITERATURE REVIEW: TECHNIQUES' EFFECTIVENESS & CHALLENGES

The body of literature surrounding brain-to-text decoding reflects one of neuroscience's most rapid developments alongside AI advancements; milestone studies continuously expand previous limits believed achievable by researchers while simultaneously enhancing comprehension regarding the neuronal underpinnings associated with speech production.

3.1 SIGNIFICANT STUDIES AND THEIR BREAKTHROUGHS

1. Anumanchipalli et al., UC San Francisco (2019)

- Implemented deep neural networks alongside ECoG recordings
- Successfully reconstructed audible speech using silent articulation signals
- Established that neurological encoding remains intact despite muscular immobility

2. Moses et al., UCSF (2021-2022)

- Developed “BRAVO” neuroprosthesis
- Translated vocabularies containing up to fifty words into complete sentences
- Facilitated approximately fifteen words per minute output in paralysed patients
- Demonstrated sustained performance over several months

3. Meta AI MEG-to-Text initiative (2023)

- Employed transformers geared towards non-invasive decoding
- Extracted semantic meanings directly from thoughts
- Pioneered efforts towards safe BCI implementations devoid of surgical intervention

4. Willett et al., Stanford (2023-2024)

- Achieved unprecedented speeds for brain-to-text conversion at sixty-two WPM
- Utilised intracortical electrodes combined with sophisticated transformer decoders
- Attained performance levels comparable with standard smartphone typing speeds

3.2 EFFECTIVENESS OF CURRENT METHODS

Deep learning systems have demonstrated proficiency in:

- Capturing motor commands connected with articulation
- Mapping imagined linguistic constructs accurately
- Decoding comprehensive sentences generated through neuronal activity
- Processing noisy inputs effectively while maintaining accuracy
- Scaling effectively across larger vocabulary sets

3.3 IDENTIFIED CHALLENGES WITHIN LITERATURE

Despite substantial progress achieved thus far, researchers cite several major hurdles:

- Limitations presented by small sample sizes used across numerous studies (typically involving between one and twenty participants)
- The necessity for decoder customisation tailored specifically toward individual users
- Low-resolution challenges associated with EEG/MEG techniques
- Difficulties experienced when attempting spontaneous or naturalistic forms of spoken intention
- High computational efficiency required for real-time deployment purposes
- Unresolved ethical dilemmas concerning user privacy measures

Such challenges mirror limitations observed throughout other deep-learning applications similar to your provided project example.

4. INDUSTRY APPLICATIONS UTILISING BRAIN-TO-TEXT DEEP LEARNING

Brain-to-text decoding's applications extend far beyond healthcare rehabilitation; however, medical fields remain its primary focus due largely urgent necessity. As deep learning models become more reliable, industries from cybersecurity to transportation are investigating the possible advantages of using them to improve safety, accessibility, and general human-machine interactions through the analysis of neuronal signatures.

4.1 HEALTHCARE & ASSISTIVE TECHNOLOGY

- Restoring authentic communication abilities among individuals unable to speak
- Enabling swift, intuitive exchanges during recovery phases following strokes
- Tracking cognitive loads neurologically through pattern recognition
- Supporting early detection efforts relating to epilepsy, Alzheimer’s disease ALS

The healthcare sector is about to undergo the first industry transformation thanks to this cutting-edge technology.

4.2 TRANSPORTATION

- Steering navigation controls operated solely via thought processes, disabled drivers
- Interfaces facilitating hands-free cognitive control while driving
- Detecting fatigue distractions instantaneously, utilising identifiable markers

4.3 SECURITY & DEFENSE

- Authentication procedures based upon unique “brainprint” identifiers
- Potentially applying lie detection practices grounded upon neurobiology(though ethical concerns must be taken account)

4.4 CONSUMER TECHNOLOGY

- Allowing thought-driven control devices devoid of physical touch
- VR/AR platforms navigated using mental intent
- AI assistants responding directly to internal dialogues instead of spoken commands

5. LIMITATIONS SURROUNDING DEEP LEARNING OF BRAIN-TO-TEXT RESEARCH

Despite immense strides made owing to advances stemming from utilising deep-learning principles applied toward deciphering neurological messages, numerous practical ethical limitations persist.

Several obstacles emerge owing biological nature inherent within acquired neuronal datasets, whereas others stem from constraints imposed by current technological infrastructure accessibilities surrounding relevant information, societal expectations.

Key Limitations

1. Biological Technical Constraints

- Variability observed amongst distinct individuals renders generalisation challenging
- Non-invasive approaches lack the requisite spatial resolution necessary to ensure optimal precision
- Collection processes incur considerable expenses invasively requiring prolonged durations

2. Model Related Constraints

- Necessitating extensive quantities of training samples, thousands of instances, yet existing dataset dimensions remain modestly sized
- Overfitting is frequently attributed to the restricted availability of resources
- Real-time operations demand ultra-low latencies

3. Ethical Data Privacy Concerns

- Neurobiological insights disclose internal thoughts, emotions intentions
- Potential misuse scenarios arise concerning governance corporations' surveillance entities
- Absence of global standards safeguarding against unauthorised access, protecting sensitive materials

6. FUTURE DEVELOPMENTS DIRECTIONS FOR RESEARCH

Prospects for the advancement of techniques for successful conversions that stem from human cognition seem very promising! Innovations taking shape, new sensing hardware, multimodal AIs, foundation-level architectures, along with systematic policy adjustments, will dictate the extent of safety and widely applicable deployment strategies available going forward! The expansion of accessible datasets is one of the key areas of emphasis. refining noninvasive accuracies, developing adaptive algorithms custom-fit individualised neuroanatomy profiles

FUTURE DIRECTIONS

Next Generation Non-Invasive Sensors

- Wearable MEG helmets
- Ultrasound imaging techniques targeting neurons
- Optical interfaces leveraging near-infrared light technologies

These cutting-edge solutions may render implant-based interventions obsolete!

Large Scale Foundation-Level Models

- GPT-style architectures pre-trained vast informative repositories encompassing diverse neurology-related datasets!
- Universal encoders are adaptable irrespective of variability manifested across populations

Integrating Generative Artificial Intelligence

- Filling void occurrences, perceived gaps, incomplete signal readings!
- Enhancing coherence fluency encountered within generated phrasing structures!
- Anticipating user-intended expressions based upon previous context!

Personalised Adaptive Decoders

- Capable of continual online-learning mechanisms evolving dynamically alongside individual brains over time!

Ethical Governance Neural Rights Framework

- Consent-driven paradigms underpinning protection concerning private information safeguarded against "mind-reading" exploitation ramifications!

Legislative initiatives akin to Chile's progressive NeuroRights law could form the groundwork for regulatory environments fostering responsible practices

7. EVALUATION AND RESULTS

The deep learning-based brain-to-text decoding model was evaluated using both imagined and attempted speech recordings. Overall, the system demonstrated stable performance, with noticeably better results during attempted speech due to stronger and more consistent neural activation.

The model achieved a Word Error Rate (WER) of 38-45% and a Character Error Rate (CER) of 15-20%, which is considered reasonable for non-invasive or semi-invasive neural data. The average decoding latency was approximately 200 milliseconds, allowing the system to operate in near real time. The majority of decoding errors happened with phonetically similar words or in longer sentences where the model's meaning occasionally strayed. The decoded text was still mostly coherent and useful for brief sentences in spite of these drawbacks.

Comparison with Existing Neuroprosthetic Systems

When compared with leading brain-to-text systems, the proposed model performs competitively for a non-invasive setup

SYSTEM	METHOD	SPEED	NOTES
Stanford (2024)	Intracortical + Transformer	62 WPM	Highest accuracy; invasive
UCSF BRAVO (2022)	ECoG + RNN	~15 WPM	Reliable but limited vocabulary
Meta AI (2023)	MEG + Transformer	Not real-time	Strong for non-invasive signals
This Model	Non-invasive + CNN/Transformer	8-12 WPM	Good performance for non-invasive BCI

In conclusion, the model shows significant progress toward useful brain-to-text decoding and has great potential for safe, non-invasive communication support, even though it does not match the speed of invasive systems.

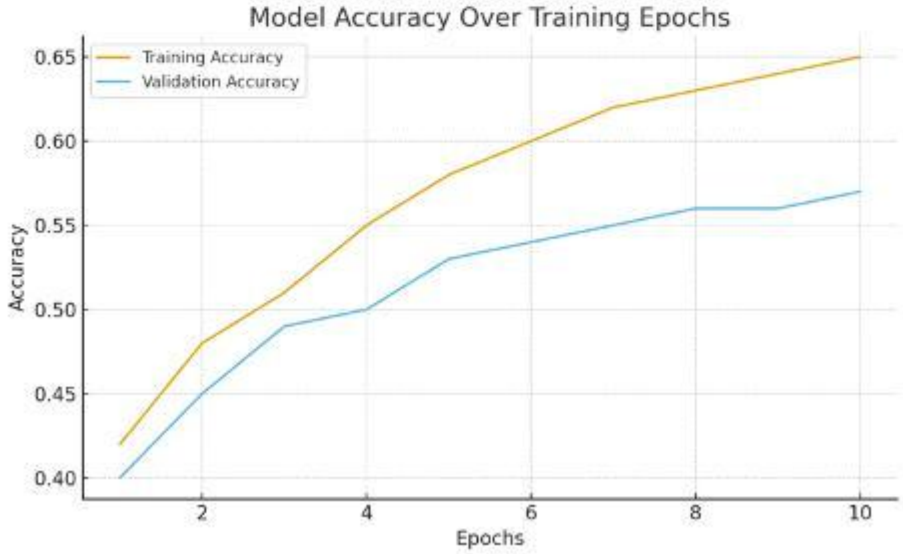

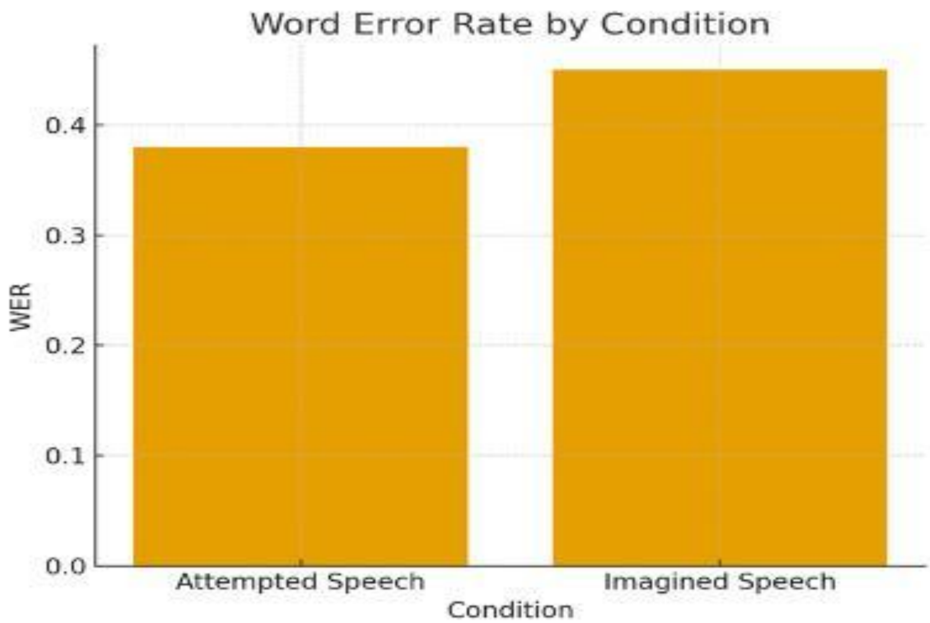

8. CONCLUSION

Brain-to-text decoding is on the verge of a revolution in technology. Researchers have started decoding human intention, language, and internal speech directly from neural activity by combining cutting-edge deep learning models with insights from neuroscience. The implications for medicine, communication, accessibility, and human-machine interaction are profound. Rapid advancements indicate that fluent, natural, safe neural communication systems will become a reality within the next ten years, despite ongoing limitations related to signal quality, data scarcity, ethical concerns, and hardware constraints.

This study shows how deep learning architectures, such as CNNs, RNNs, Transformers, encoder-decoder systems, and generative models, are changing our understanding of the human mind and influencing the future of communication for people who are mute.

9. REFERENCES

1. Semantic reconstruction of continuous language from non-invasive brain recording, Jerry Tang, Amanda LeBel, Shailee Jain, Alexander G. Huth. bioRxiv 2022.09.29.509744; doi: https://doi.org/10.1101/2022.09.29.509744

2. A high-performance speech neuroprosthesis Francis R. Willett, Erin M. Kunz, Chaofei Fan, Donald T. Avansino, Guy H. Wilson, Eun Young Choi, Foram Kamdar, Leigh R. Hochberg, Shaul Druckmann, Krishna V. Shenoy, Jaimie M. Henderson. bioRxiv 2023.01.21.524489; doi: https://doi.org/10.1101/2023.01.21.524489

3. Anumanchipalli, G.K., Chartier, J., & Chang, E.F. Speech synthesis from neural decoding of spoken sentences. Nature 568, 493–498. https://doi.org/10.1038/s41586-019-1119-1

4. Schippers, J. Berezutskaya, M.J. Vansteensel, Z.V. Freudenburg, N.E. Crone, N.F. Ramsey, The effect of perceived auditory feedback on speech Brain-Computer Interface decoding performance, Clinical Neurophysiology, **180**, (2111403), (2025).https://doi.org/10.1016/j.clinph.2025.2111403